

[Forums / Module 5 forum](#)[Help](#)

## problem with VM - tried to start it on Virtual Box - SUSE Linux 12.3 host (64 bit)

 You are subscribed. [Unsubscribe](#) [virtualMachineError](#) × [+ Add Tag](#)Sort replies by: [Oldest first](#) [Newest first](#) [Most popular](#)[Karen West](#) · 11 days ago 

I am running SUSE Linux 12.3 64 bit Intel CPU as host. I downloaded your spatial\_computing.vdi. I downloaded and installed Virtual Box and tried to start this VM (spatial\_computing.vdi) and I received the following error message. If you can help solve this issue, please let me know.

Otherwise, a few weeks ago, I installed all the recommended software directly on my 9 year old daughter's netbook that runs Ubuntu Linux on it, although I was not sure what I was doing when I installed all that from the links provided at the start of this class, before I had heard about you providing it all in a VM. I prefer to work in the VM if you know how to fix this, since I can be confident that you installed all the correct software, and not worry that when I did it, I did some wrong version somewhere. However, the reason I may use the netbook instead, is because I once started another class where the VM was provided, and I received the same error in that class, and had to withdraw at the time because I could not get the VM to run on this SUSE Linux machine I have, and in that class, there was no alternative. So please let me know if you know what is wrong here. In the class I took one year ago - no one on the discussion forums, nor the teaching assistants knew what was wrong, since no one was using this version of SUSE Linux as host. So let me know if you do, or otherwise I will attempt it using the netbook instead, and if that has issues - well I'll just skip the project part of this course, and just do concepts instead. Thanks.

Here is the error message:

The VirtualBox kernel modules do not match this version of VirtualBox. The installation of VirtualBox was apparently not successful. Executing

```
'/etc/init.d/vboxdrv setup'
```

may correct this. Make sure that you do not mix the OSE version and the PUEL version of VirtualBox.

↑ 0 ↓ · flag

[Toby Jia-Jun Li](#) STAFF · 11 days ago 

I suspect you have another version of VirtualBox (or the kernel) sitting somewhere in the machine...Have you

tried Google the error message you got?

Working on that netbook is definitely possible. You can either (1) Try to get the VM running on the netbook or (2) Follow the instruction on the WikiBrain homepage to load Simple English, Wikidata and spatial data and then get the sample program from the coursera page.

↑ 0 ↓ · flag



Karen West · 11 days ago



Thank you for the response, Toby - here is a long cut-and-paste type question--and I apologize in advance--I'm not sure if I can complete the projects for this class, so I'm trying to figure that out.

A year ago when that error occurred I did google the error, tried to inquire in various places, found no help, and I withdrew from that class actually. (It was part 2 of the Agile Web Apps Ruby on Rails actually - so I only completed part 1 since I fell behind trying to figure out this installation error with the VM and had to withdraw). If someone else can help me get this question answered, perhaps I'll use it, but I found last time trying to solve it resulted from withdrawing from the class. This class I know if I cannot get the tools working I can complete the "Concepts" part without doing the project, but we'll see how it goes.

So thanks Toby - I think I'm going to try to do step 2 - and use Eclipse which was recommended I think when I installed everything on the netbook. I'm a bit lost though with the importing directions and wondering if you can help here. I've only used Eclipse once, in the Android Apps class I took, so sorry for my ignorance here. That was my only exposure to the Java and XML languages to complete a program that uses them. (I used to be an embedded C programmer for 18 years before I lost my job just over 5 years ago, and a few years ago, started to expand online learning new programming languages -- I come from an EE education, way back when. ;-)). So I apologize in advance what others probably already know in this class!

I installed the following directories on the netbook with Ubuntu, in addition to Eclipse Java EE IDE:

geos-3.3.9  
h2database  
maven  
postgis-2.0.6  
wikibrain

This was during the first week of the class, in the announcements post that provided these links.

Here are my specific questions on the trouble I'm having with eclipse to import wikibrain (you mentioned above--Simple English, Wikidata and spatial data)--and try as I may with the wikibrain homepage import instructions --I'm lost! Please see my QUESTIONS in capitals below. Please let me know which import method would be best for me, given the above explanation.

## Importing data

The Graphical Loader offers the easiest way to import data. However, WikiBrain supports much

more sophisticated importing configurations.

## GraphicalLoader

You can launch the graphical loader by double-clicking the wikibrain-withdeps jarfile, or by running the `org.wikibrain.GuiLoader` class.

QUESTION--where is this graphical loader? (`org.wikibrain.GuiLoader`)

QUESTION--is the wikibrain-withdeps jarfile in the wikibrain directory I downloaded from the link provided the first week of this class? If I find it, where do I double click it?

Here is more cut and paste from the import page on the wikibrain page: Please see my QUESTIONS.

You can customize basic features of the loading process using the graphical loader. Be aware that the loader writes a configuration file named `customized.conf` that you'll need to use when running your java programs (i.e. `-c customized.conf`).

QUESTION: Where is customized.conf?

## Customized loading setups

You can achieve more customized loader setups by running the Loader program directly. You can do this your IDE's run dialog, or by downloading and installing the `wb-java.sh` script described in the Installation page.

WikiBrain's Loader class manages the downloading, parsing, and import of [Wikipedia datasets](#) published by the Wikimedia foundation.

QUESTION: Where is `org.wikibrain.Loader`? How do I run it?

QUESTION: You mentioned importing Simple English, and below it says that the loader runs without any arguments except a language, which downloads, parses and imports the latest database files for the Simple English language edition of Wikipedia into an embedded h2 database. I know I also downloaded above a directory called "h2database". How is this importing Simple English into that database from Eclipse? I am lost here again, but I did download all the required directories mentioned above.

Most commonly, Loader is run without any arguments except a language. This downloads, parses and imports the latest database files for the Simple English language edition of Wikipedia into an embedded h2 database:

```
org.wikibrain.Loader -l simple
```

Each step in the WikiBrain loading process is called a **stage**. By default a set of `core` stages are loaded (articles, links, lucene search index), but others are not (wikidata, spatial). You can specify the stages you want loaded, such as wikidata using the `-s` option:

```
org.wikibrain.Loader -l simple -s wikidata
```

You may also want to configure how the stages you select are run. For example, you can specify that each stage should incorporate a custom override configuration file:

```
org.wikibrain.Loader -l simple -s wikidata -c custom.conf
```

These features are explained in more detail below.

## Stage selection

WikiBrain's import pipeline is divided into stages, listed in the table below. If you don't tell WikiBrain to load specific stages, the stages marked **on = Y** are run. The **concepts** stage is also run by default if more than one language is installed.

on	stage	dependsOn	description
Y	fetchlinks	nothing	Downloads a list of available data files in the requested languages.
Y	download	fetchlinks	Downloads any new data files in the requested languages.
Y	dumploader	download	Loads articles and raw pages from the dump files.
Y	redirects	dumploader	Loads redirect mapping.
Y	wikitext	redirects	Loads links and categories.
Y	lucene	wikitext	Builds lucene full-text search indexes.
Y	phrases	wikitext	Builds and loads phrase to page disambiguation mappings.
?	concepts	redirects	Builds and loads mapping from language-specific articles to universal concepts.
N	universal	concepts, wikitext	Builds "universal" links.
N	wikidata	concepts	Downloads and loads wikidata facts for the requested languages.
N	spatial	wikidata	Installs geospatial data.
N	sr	wikitext, phrases, lucene	Builds semantic relatedness models.

You can run a specific stage using the **-s** flag. For example, the following command builds the semantic relatedness (SR) model and prepares it for use.

```
org.wikibrain.Loader -s sr
```

You can specify multiple stages using multiple **-s** arguments:

```
org.wikibrain.Loader -s sr -s spatial
```

If you ask WikiBrain to run a stage (e.g. **sr**), it checks to see if the stages it depends on (in this case, **wikitext**, **phrases**, **lucene**) have been run. If necessary it runs them, and any further transitive dependencies. You can disable checking to see if a stage is run, and rerun all necessary stages by using the **-d** dependency (for drop).

## Configuring stages

You can configure the global settings for all stages by passing any "standard" WikiBrain argument. If any of these options exist, they will be passed along to all stages.

<code>-c, --conf</code>	configuration file
<code>-h, --threads</code>	the maximum number of threads that should be used
<code>-l, --languages</code>	the set of languages to process, separated by commas or 'LOADED'
<code>--base-dir</code>	the base directory used to resolve relative directories
<code>--tmp-dir</code>	the temporary directory

In addition, you can pass stage-specific arguments via the `-s` flag. For example the following command runs the phrase stage (turns it `on`), and passes the argument `-p stanford` to it to tell it to load the [Stanford concept analyzer](#).

```
org.wikibrain.Loader -s "phrases:on:-p stanford"
```

## More details

Each WikiBrain stage corresponds to a Java class with a `main()` method that can be run as a standalone Java program. For example, the wikitext stages corresponds to [org.wikibrain.loader.WikiTextLoader](#). You can find the full definition of stages at the end of the [reference.conf](#).

↑ 1 ↓ · flag

[+ Comment](#)



Toby Jia-Jun Li STAFF · 11 days ago

- Q1. GUI loader is the one pop up when you double-click the WikiBrain jar file (You can also run the jar file from the command line in the case your system does not let you run a java executable from File Manager)
- Q2. It's a single jar file that can be found at [www.wikibrainapi.org](http://www.wikibrainapi.org), on the right panel I remember.
- Q3. If you decide to use the GUI loader, there is no need to use the Loader class

Before loading the data, you should connect to your postgresSQL database, create a new database and name it whatever you like("wikibrain" let's say).

Then you should enable PostGIS in that database

Back to the GUI loader, for language, you should fill in "simple". Select "SR", "Wikidata" and "Spatial" in

addition to the default selections. Database type should be postgres. Also fill in your postgres database name, your postgres username and password.

↑ 2 ↓ · flag



Karen West · 11 days ago



Thank you Toby - I will give this a try again later and let you know if I get it working, or if I get stuck and have more questions. I have to take a short break from it at the moment.

↑ 0 ↓ · flag



Karen West · 7 days ago



I'm just getting back to this today. The end of last week was disrupted by doctor visits and an emergency MRI for my husband's head after injuring it, so I'm still recovering to the schedule disruption. I might try first upgrading my SUSE Linux, and if that does not help with Oracle's Virtual Box loading the Ubuntu VM, then I'll continue with the advice above for the Ubuntu netbook that I have. The reason it makes sense to first try upgrading SUSE Linux, is because I had also started the Software Security class last week, and they too use Oracle's Virtual Box with a Ubuntu VM, and they said I should try upgrading my SUSE Linux from 12.3 to 13.1, and that in doing so I will only disrupt my Linux system files and nothing else on my disk. So I'll try that and I already have the spatial computing VM, and if that does not help, well then, the last resort will be to try to get it going on the netbook, or otherwise just stick with the CONCEPTS version of this class, rather than the TECHNICAL TRACK. Thanks for the help last week.

↑ 0 ↓ · flag



Karen West · 7 days ago



I have just rebooted after saving my entire disk before upgrading to SUSE Linux 13.1. Tomorrow I will try Oracle VirtualBox again!

↑ 0 ↓ · flag



Karen West · 6 days ago



Good morning. I uninstalled Oracle Virtual Box from the new SUSE Linux 13.1 version, and I just re-installed it. I just re-tried to start Oracle Virtual Box spatial computing Ubuntu VM for this class and I'm very disappointed to report that after all that effort to save my entire disk yesterday, followed by taking quite awhile to upgrade my SUSE Linux from 12.3 to 13.1 version - re-starting the spatial computing Ubuntu VM gets the exact same error as before that I mentioned in this message chain. It's as if what ever was broken appeared to uninstall--but it must not have, since the error is still there. Unless you know the answer to why the VM won't work for me after upgrading Linux yesterday - the only other option for me is to once again attempt to get this

working on my Ubuntu netbook that does not use VirtualBox or the VM. Otherwise I'll just do the CONCEPTS version of this class since it will be impossible to attempt your projects of course with no tools!

↑ 0 ↓ · flag

+ Comment



Karen West · 6 days ago



In the below mentioned advice - how do you "connect to your postgresSQL database, create a new database and name it whatever you like("wikibrain" let's say). Then you should enable PostGIS in that database" ?? Is this from Eclipse? Where do you connect to the postgresSQL database? Sorry to be ignorant here but I've never done this before.

I just downloaded the "wikibrain-withdep;s-0.4.1.jar" file - this was what I found as you mentioned in the advice "in the right panel of [www.wikibrainapi.org](http://www.wikibrainapi.org)" - where you said I should AFTER doing the above step of "connecting my postgresSQL database and creating a new wikibrain database" - I should double click on the Wikibrain .jar file - is this the correct file I saw in the right panel where I then load the data?

I was not sure you meant in the response you gave, but given I wanted to first try to upgrade my SUSE Linux from 12.3 to 13.1, and then re-install Oracle's Virtual Box and attempt to start your spatial computing Ubuntu Linux VM instead of having to load all this stuff myself on my Ubuntu netbook - and now that has failed - I am just asking questions on this advice you gave last Friday for installing on the Ubuntu Linux netbook.

Please clarify the above questions - thank you.

I am recopying the advice you gave above on Friday below.

Q1. GUI loader is the one pop up when you double-click the WikiBrain jar file (You can also run the jar file from the command line in the case your system does not let you run a java executable from File Manager

Q2. It's a single jar file that can be found at [www.wikibrainapi.org](http://www.wikibrainapi.org), on the right panel I remember.

Q3. If you decide to use the GUI loader, there is no need to use the Loader class

Before loading the data, you should connect to your postgresSQL database, create a new database and name it whatever you like("wikibrain" let's say).

Then you should enable PostGIS in that database

Back to the GUI loader, for language, you should fill in "simple". Select "SR", "Wikidata" and "Spatial" in addition to the default selections. Database type should be postgres. Also fill in your postgres database name, your postgres username and password.

↑ 0 ↓ · flag



Toby Jia-Jun Li STAFF · 6 days ago 🔗

Yep that's the file.

You should use psql to connect to your postgres database and create a database first.

If your postgres installation is completed, type psql in terminal should connects you to the database.

↑ 0 ↓ · flag



Karen West · 6 days ago 🔗



I'm apologize in advance - but I've never done this before - what do you mean by "use psql to connect to my postgres database" - from eclipse - or from the file manager or a terminal console?

↑ 0 ↓ · flag



Toby Jia-Jun Li STAFF · 6 days ago 🔗

From terminal, type "psql" and press enter. You should see the postgres database if your postgres is installed properly. Otherwise you'll need to install postgres first.

↑ 0 ↓ · flag



Karen West · 6 days ago 🔗



Hi Toby - I do not have my postgres database installed correctly - the output of psql was:

```
sophie@dwarf:~/karen$ psql
```

```
psql: FATAL: role "sophie" does not exist
```

The link to the instructions I followed that first week is:

<http://trac.osgeo.org/postgis/wiki/UsersWikiPostGIS20Ubuntu1204src>

Please let me know if you can further recommend where I've gone wrong! I could not by the way get another class's Ubuntu VM to work either - seems to be some issue with something called DKMS package, even after upgrading SUSE Linux from 12.3 to 13.1 yesterday.

Here is what I downloaded - all the directories where there were links posted the first week of class - I had no idea what I was downloading and installing that day since this is all new to me! I'm not sure what I did wrong.

```
sophie@dwarf:~/karen$ ls postgis-2.0.6
aclocal.m4  configure.ac  libpgcommon  postgis_config.h.in
astyle.sh   COPYING      libtool      postgis_svn_revision.h
authors.svn CREDITS      LICENSE.TXT  raster
autogen.sh doc          loader       README.postgis
ChangeLog  extensions  ltmain.sh    regress
config.guess  extras      macros       spatial_ref_sys.sql
```



```

config.log  GNUmakefile  Makefile      STYLE
config.rpath GNUmakefile.in  MIGRATION  TODO
config.status  install-sh  NEWS      topology
config.sub    java      postgis    utils
configure     liblwgeom   postgis_config.h  Version.config
sophie@dwarf:~/karen$

```

I'll list the other directories I downloaded the first week of this class on my daughter's netbook, since they too are related to databases and I don't know if I was supposed to do something with them (remember I'm EE educated - did 18 years in embedded C, and recently expanded in CS online - I have taken a databases class, but the databases we interfaced our code too were all elsewhere so I've never set them up-so apologies!):

```

sophie@dwarf:~/karen$ ls h2database
h2
sophie@dwarf:~/karen$ ls h2database/h2
bin  build.bat  build.sh  docs  service  src
sophie@dwarf:~/karen$

```

```

sophie@dwarf:~/karen$ ls wikibrain
antlr          ill_langs.csv      net
assets         javassist          org
au             javax              osmosis-plugins.conf
bonecp-default-config.xml  jnt                README
cern           joptsimple         readme.txt
ChangeLog      jsr166y            reference.conf
com            junit              Skunkworks.class
commons-logging.properties  language_info2.tsv  spatial-matcher
corejava       language_info.tsv  stopwords
datasets       languages.tsv      templates
db             LICENSE            testpool.jocl
de             license.html       unicode-license.txt
dictionary.py   licenses           wb-java.sh
edu            license.txt        xjb
EDU            LICENSE.txt        XMLPULL_1_1_3_1_VERSION
fooz.txt       log4j.xml          XPP3_1.1.4c_MIN_VERSION
gnu            make_lang_constants.py  xsd
hep            META-INF           xtc

```

```

sophie@dwarf:~/karen$ ls geos-3.3.9
aclocal.m4  CMakeLists.txt  COPYING  ltmain.sh  NEWS  tests
acsite.m4  config.guess  depcomp  macros  nmake.opt  TODO
AUTHORS  config.log  doc  Makefile  php  tools
autogen.bat  config.status  include  Makefile.am  py-compile
capi  config.sub  INSTALL  Makefile.in  README
ChangeLog  configure  install-sh  makefile.vc  src
cmake  configure.in  libtool  missing  swig

```

```
sophie@dwarf:~/karen$ ls maven
apache-maven-3.2.3
sophie@dwarf:~/karen$ ls maven/apache-maven-3.2.3/
bin boot conf lib LICENSE NOTICE README.txt
sophie@dwarf:~/karen$

sophie@dwarf:~/karen$ ls install-tilemill.sh
install-tilemill.sh

sophie@dwarf:~/karen$ ls eclipse
eclipse
sophie@dwarf:~/karen$ ls eclipse/eclipse/
about_files configuration eclipse.ini icon.xpm plugins
about.html dropins epl-v10.html notice.html readme
artifacts.xml eclipse features p2
sophie@dwarf:~/karen$

sophie@dwarf:~/karen$ ls workspace/
Module 5 programming assignment.pdf RemoteSystemsTempFiles
QueryExample.java WikiBrain-WikiSym2014.pdf
QueryExample_Project1 WikiBrain_WikiSym.pdf
sophie@dwarf:~/karen$
```

↑ 0 ↓ · flag

[+ Comment](#)



**Toby Jia-Jun Li** STAFF · 6 days ago 🔔

Nice! That means postgres is already there...just no user was created.

Do "sudo su - postgres" to switch to the system user "postgres", and then do "psql"

After logged into psql, create a new user. And then log in with the new user to create new database for WikiBrain and enable PostGIS in the new database.

↑ 0 ↓ · flag



Karen West · 6 days ago



Before I take the next step I thought I should double check with you: here is the output of `sudo su - postgres`. Are there instructions somewhere on how to create a new user, login with new user to create new database for Wikibrain and enable PostGIS in the new database?

```
sophie@dwarf:~/karen$ sudo su - postgres
sudo: unable to resolve host dwarf
[sudo] password for sophie:
Sorry, try again.
[sudo] password for sophie:
postgres@dwarf:~$ psql
psql (9.1.14)
Type "help" for help.
```

```
postgres=#
```

↑ 0 ↓ · flag



Toby Jia-Jun Li STAFF · 6 days ago

Yep you are now at the right place!

Creating new user / new database should be easily Google-able. Then you connect to your new database and run "CREATE EXTENSION postgres;" to enable PostGIS.

You should then fill in your username (probably "sophie") and the name of your newly created database when running the WikiBrain setup GUI.

↑ 0 ↓ · flag



Karen West · 6 days ago



Toby - I just googled the following for CREATE DATABASE for postgres - can you help with the parameters I should supply?

# CREATE DATABASE

## Name

CREATE DATABASE -- create a new database

## Synopsis

```
CREATE DATABASE name
```

```
[ [ WITH ] [ OWNER [=] user_name ]  
  [ TEMPLATE [=] template ]  
  [ ENCODING [=] encoding ]  
  [ LC_COLLATE [=] lc_collate ]  
  [ LC_CTYPE [=] lc_ctype ]  
  [ TABLESPACE [=] tablespace ]  
  [ CONNECTION LIMIT [=] connlimit ] ]
```

## Description

CREATE DATABASE creates a new PostgreSQL database.

To create a database, you must be a superuser or have the special CREATEDB privilege. See [CREATE USER](#).

By default, the new database will be created by cloning the standard system database template1. A different template can be specified by writing TEMPLATE name. In particular, by writing TEMPLATE template0, you can create a virgin database containing only the standard objects predefined by your version of PostgreSQL. This is useful if you wish to avoid copying any installation-local objects that might have been added to template1.

## Parameters

name

The name of a database to create.

user\_name

The role name of the user who will own the new database, or DEFAULT to use the default (namely, the user executing the command). To create a database owned by another role, you must be a direct or indirect member of that role, or be a superuser.

template

The name of the template from which to create the new database, or DEFAULT to use the default template (template1).

encoding

Character set encoding to use in the new database. Specify a string constant (e.g., 'SQL\_ASCII'), or an integer encoding number, or DEFAULT to use the default encoding (namely, the encoding of the template database). The character sets supported by the PostgreSQL server are described in [Section 22.2.1](#). See below for additional restrictions.

lc\_collate

Collation order (LC\_COLLATE) to use in the new database. This affects the sort order applied to strings, e.g. in queries with ORDER BY, as well as the order used in indexes on text columns. The default is to use the collation order of the template database. See below for additional restrictions.

lc\_ctype

Character classification (LC\_CTYPE) to use in the new database. This affects the categorization of characters, e.g. lower, upper and digit. The default is to use the character classification of the template database. See below for additional restrictions.

tablespace

The name of the tablespace that will be associated with the new database, or DEFAULT to use the template database's tablespace. This tablespace will be the default tablespace used for objects created in this database. See [CREATE TABLESPACE](#) for more information.

connlimit

How many concurrent connections can be made to this database. -1 (the default) means no limit.

Optional parameters can be written in any order, not only the order illustrated above.

↑ 0 ↓ · flag



Karen West · 6 days ago 🔗



Toby - realistically - with the delays that seem to be happening with trying to get the tools installed to do your assignment for the technical track - unless I get some specific help rather quickly with setting it up on the netbook, most likely I will have to just do the concepts track. Please make a recommendation. If the set up is not too bad, if you could give more specific steps on how to set it up there, given I gave you all the directories and files I downloaded - I don't think I'm going to just guess on how to set it up by doing web searching - with out advising from you or someone else in the class. It's beginning to sound like a time waste, unless I can get the tools installed and do the actual assignment that was assigned to this class - which was not the tool installation. So please give your advice here. Thank you.

↑ 0 ↓ · flag



Toby Jia-Jun Li STAFF · 6 days ago 🔗

Will actually I'm not supposed to provide individual technical support on setting up, as stated in the syllabus. If you choose to not use the VM, you are supposed to be on your own...

Getting back to your question, you won't need any additional parameter other than the name of the database.

↑ 0 ↓ · flag



Karen West · 6 days ago 🔗



Yes - I realize that - thank you for your response. Maybe I can get through it, who knows. I wonder why no one else is having these troubles with the VM if they too have SUSE Linux 13.1? I'm having the same trouble in another class as well - their Ubuntu VM does not work either - and in that class, withdrawal is the only option even though I did well on the first few quizzes, since they do not have the "concepts only" option to it, nor is there an option to do their projects on the netbook with Ubuntu. I have not yet given up yet, but we'll see how it goes. It must seem a bit funny to you too though, that I'm the only one experiencing this much trouble with getting the Ubuntu VMs to work on my SUSE Linux 13.1 just upgraded yesterday, or perhaps I'm the only one using that OS? Anyway, thank you for your response to help.

↑ 0 ↓ · flag

[+ Comment](#)

Karen West · 6 days ago



Thank you Toby for trying to help me solve my issues to attempt the technical track for this class. Today I worked with some people in another class to try to make Oracle Virtual Box on my newly upgraded SUSE Linux. Given it still is not working, and one of the students in the class even took control of my computer using a tool I had not heard of called teamviewer, and I tried many different people's advice on how to get this to work for me - I tried every option recommended, and it continues not to work. I guess these things happen for a reason. There's just not enough time to also try to get it all working on my own on the netbook when I've never done that before, although I'm sure the installation is straight forward if you have. I'll decide for sure tomorrow, but I guess there is no choice but to either just complete the concepts part of the course, or withdraw from the course at this point as I will have to do with the other course. But I wanted to thank you for trying to help, and I'll decide for sure tomorrow.

↑ 0 ↓ · flag

[+ Comment](#)

Karen West · 4 days ago



So my last attempt is to download VMware's VM stuff and see if that works - however - I have to download a trial version (no money!) and I also do not know what I'm doing - so if anyone can advise here on which one I download, please let me know! Here is what I found (for desk top and end-user computing list - not the other lists of versions of their tools): If anyone knows which trial version I should download, please let me know! Time is not up yet - it's Friday and my kids are celebrating Halloween tonight - but who knows - I may pull it off - since I do know C, some OS and architecture stuff fairly well - it's the VM tools that I'm lost on! ;-)

#### Desktop & End-User Computing

- VMware Horizon (with View)  
[Download ProductDrivers & ToolsDownload Trial](#)
- VMware Horizon DaaS  
[Download ProductDrivers & Tools](#)
- VMware Horizon View  
[Download ProductDrivers & ToolsDownload Trial](#)
- VMware Horizon Clients  
[Download ProductDrivers & Tools](#)
- VMware CloudVolumes  
[Download ProductDrivers & Tools](#)
- VMware Workspace  
[Download ProductDrivers & ToolsDownload Trial](#)
- VMware Mirage  
[Download ProductDrivers & ToolsDownload Trial](#)
- VMware vCenter Operations Manager for Horizon  
[Download ProductDrivers & ToolsDownload Trial](#)
- VMware vCenter Operations Manager for Horizon View  
[Download ProductDrivers & ToolsDownload Trial](#)
- VMware ThinApp

[Download ProductDrivers & Tools](#)

- VMware Workstation

[Download ProductDrivers & ToolsDownload Trial](#)

- VMware Fusion

[Download ProductDrivers & ToolsDownload Trial](#)

- VMware Player

[Download Product](#)

- Other

- See more at: <https://my.vmware.com/web/vmware/downloads#sthash.jVOF58l.dpuf>

↑ 0 ↓ · flag



Karen West · 4 days ago 🔗



oops -cut and pasted that from another class - so ignore the last part about skills for the project!  
that class uses the VM too.

↑ 0 ↓ · flag



Karen West · 4 days ago 🔗



VMware's Player does not read in .vdi VM files? It reads in .ova files. The spatial computing VM is .vdi. Oh well - I did try once again to see if another tool would work for me to attempt your programming assignment! ;-) It looks like concepts is the only option, if I can get that done on time too now (put all my time this week for this class trying to make the VM tools work!) So we'll see. I did try to attempt your project though once again. Happy Halloween!

↑ 0 ↓ · flag

[+ Comment](#)



Karen West · 3 days ago 🔗



Someone in another class recommended the following - that class uses a Ubuntu Linux mooc-vm.ova image. He said:

---

I didn't want to install VirtualBox just for this class, so I opted to use KVM/QEMU. I was able to get the correct answers for the first project running the provided image in this manner.

```
tar -xvf mooc-vm2.ova
qemu-img convert -f vmdk -O raw mooc-vm-disk1.vmdk mooc-vm-disk1.img
qemu-system-x86_64 -enable-kvm -m 8G mooc-vm-disk1.img
```

Might be worth a shot if nothing else has worked yet.

---

So I tried it and found:

I installed KVM QEMU today and started to convert the mooc-vm to give it a try - and I ran into trouble again - here is the output - do you know what the problem is here (I know you said it worked for you):

```
scary:/home/karen/dev/softwareSecurity # tar -xvf mooc-vm.ova
mooc-vm.ovf
mooc-vm-disk2.vmdk
```

```
scary:/home/karen/dev/softwareSecurity # qemu-img convert -f vmdk -O raw mooc-vm-disk2.vmdk
mooc-vm-disk2.img
qemu-img: 'image' uses a vmdk feature which is not supported by this qemu version: VMDK version 3
qemu-img: Could not open 'mooc-vm-disk2.vmdk': Wrong medium type
qemu-img: Could not open 'mooc-vm-disk2.vmdk'
```

So do you know if KVM QEMU will work with .vdi files? In this class in which I was supposed to use VirtualBox, a Ubuntu Linux VM with a .vdi extension was provided instead of .ova, and VirtualBox did not like either of these Ubuntu VMs when I installed them. So I was wondering if you knew what I could try with the commands you provided to convert the .vdi image to work with KVM QEMU. Thanks - still trying here and there to try to get this to work, so I can try the project.

↑ 0 ↓ · flag

[+ Comment](#)



Karen West · 2 days ago 🔒



Thank you all for your help. I just tried yet another way of installing Oracle Virtual Box suggested to me:  
<http://www.unixmen.com/install-virtualbox-4-3-opensuse-13-1/>

My husband has designed his own virtualized operating system for this machine and he says it works perfectly fine on this machine - but it's just a research project - but it does so far run Linux on it - it's called Quest-V and so far it can run Linux as a VM - so he told me try his since he knows that works here - but it's just his research group's project for his Ph.D. students (which I am not and never will be!) - if something went wrong I would not know what's going on. It's not a released product - but he said if I really want to do these online courses that require VMs that just won't work on my SUSE Linux box that perhaps I should consider running Quest-V instead.

I think I've tried everything suggested - so with the deadline for the technical track project 1 being tomorrow, the concepts track will be the only option for me for this class, unless I want to run Quest-V on this machine! It's only ever been used by his research group so I admit it's a bit intimidating for me to consider it.

So here is my system set up that keeps rejecting all the latest VM distributions out there. I've tried Oracle Virtual Box, VMware Player (and my husband worked there one year too while on sabbatical, designed their caching system, so I know he has used that tool that year too), and yesterday I also tried KVM QEMU. If there are any suggestions left out there please let me know. I don't think I can use Quest-V. I have my daughter's netbook that has Ubuntu Linux on it but I do not have confidence that I would know how to set that up, and I know the people for this class are supporting many online students from many different backgrounds, so the best option for graders is the VM set up so that everyone is using the same thing.



```
karen@scary:~> lscpu
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Byte Order:        Little Endian
CPU(s):            4
On-line CPU(s) list: 0-3
Thread(s) per core: 2
Core(s) per socket: 2
Socket(s):         1
NUMA node(s):      1
Vendor ID:         GenuineIntel
CPU family:        6
Model:            42
Model name:        Intel(R) Core(TM) i3-2100 CPU @ 3.10GHz
Stepping:          7
CPU MHz:           3069.000
BogoMIPS:          6185.69
Virtualization:    VT-x
L1d cache:         32K
L1i cache:         32K
L2 cache:          256K
L3 cache:          3072K
NUMA node0 CPU(s): 0-3

karen@scary:~> uname -r
3.11.10-21-desktop

karen@scary:~> cat /etc/issue
Welcome to openSUSE 13.1 "Bottle" - Kernel \r (\l).
```

↑ 0 ↓ · flag

[+ Comment](#)



**Karen West** · 16 hours ago 🔔



Wow - with help from some people in another class - I was able to get the KVM QEMU running a Ubuntu VM running on my SUSE Linux 13.1 machine. Is there any chance that will work for this class too? I had to convert a .ova VM file to .img file instead, and that had some issues, but with help, we got it to work today. They extended their class project a week because of the troubles, and I suspect that can't be done in this class, but perhaps I will attempt your project after the due date (project 1) to see how it works (or run the solution) if the spatial\_computing.vdi can be converted to run on KVM QEMU. In that way, perhaps I can learn from the experience without having done project 1 (and similarly for project 2), if it's possible. I realize it's too late now, since I just got the VM tools working today! I sometimes like to "learn from the experience" though, if it's possible. I did take the concepts quiz that was due today, but was unable to attempt the project. So please let me know. Here are the commands that worked to convert the .ova to .img - will it work for converting .vdi to .img?

posted by someone else to try:

```
tar -xvf mooc-vm2.ova
qemu-img convert -f vmdk -O raw mooc-vm-disk1.vmdk mooc-vm-disk1.img
qemu-system-x86_64 -enable-kvm -m 8G mooc-vm-disk1.img
```

what happened for me:

```
scary:/home/karen/dev/softwareSecurity # tar -xvf mooc-vm.ova
mooc-vm.ovf
mooc-vm-disk2.vmdk
```

```
scary:/home/karen/dev/softwareSecurity # qemu-img convert -f vmdk -O raw mooc-vm-disk2.vmdk
mooc-vm-disk2.img
qemu-img: 'image' uses a vmdk feature which is not supported by this qemu version: VMDK version 3
qemu-img: Could not open 'mooc-vm-disk2.vmdk': Wrong medium type
qemu-img: Could not open 'mooc-vm-disk2.vmdk'
```

the recommendation that helped:

You already mentioned that qemu-kvm and qemu-img return error "'image' uses a vmdk feature which is not supported by this qemu version: VMDK version 3". Sometimes it is possible to change 2 bytes in file and overpass checking the version of vmdk. This method described in article <http://blog.carlos-spitzer.com/2013/12/26/how-to-convert-vmdk-3-images-to-qcow2-even-when-qemu-img-f...>

Use script from article to modify vmdk image

Then Convert vmdk to raw with qemu-img

```
qemu-img convert -f vmdk -O raw mooc-vm-disk2.vmdk mooc-vm-disk2.img
```

Here is what happened:

```
scary:/home/karen/dev/softwareSecurity # ./vmdk3-hack.sh mooc-vm-
disk2.vmdk
mooc-vm-disk2.vmdk is VMDK3.n Patching to VMDK2.n Run this script again when you're done to patch it
back.
```

```
scary:/home/karen/dev/softwareSecurity # qemu-img convert -f vmdk -O raw mooc-vm-disk2.vmdk
mooc-vm-disk2.img
```

```
scary:/home/karen/dev/softwareSecurity # ls
courseInfo  mooc-vm      mooc-vm-disk2.vmdk  mooc-vm.ovf  week0
.directory mooc-vm-disk2.img mooc-vm.ova        vmdk3-hack.sh week1
```

So you can see that running your script you recommended helped and I got beyond the qemu-img convert. However - it fails on the next thing Andrew recommended above - do you know how to fix this (allocation issue)? Thank you for your help.

```
scary:/home/karen/dev/softwareSecurity # qemu-system-x86_64 -enable-kvm -m 8G mooc-vm-disk2.img
Warning: default mac address being used, creating potential for address conflict
Failed to allocate 8589934592 B: Cannot allocate memory
```

Aborted

The final thing that made it work:

According to the QEMU docs, the `-m`` option specifies the virtual RAM for the VM. From the error message it looks like you don't have 8G of free memory available (and you don't need that much for the provided VM). Try using just 1G.

(or what ever memory would be required for the spatial\_computing.img file if it can be converted from .vdi to .img.)

↑ 0 ↓ · flag

[+ Comment](#)

New post

To ensure a positive and productive discussion, please read our [forum posting policies](#) before posting.

<b>B</b>	<i>I</i>			Link	<code>&lt;code&gt;</code>	Pic	Math		Edit: Rich ▼	Preview
<div></div>										

☐ Make this post anonymous to other students

☒ Subscribe to this thread at the same time

Add post

problem with VM - tried to start it on Virtual...

<https://class.coursera.org/spatialcomputing-0...>